

# Speaker Identification System

---

## Project Background

### Introduction

Speaker recognition is the process of **automatically identifying who is speaking** based on unique characteristics embedded in speech signals. This technology is widely used in applications such as **biometric authentication, security systems, and voice assistants**.

This project builds a **simple yet effective speaker recognition system** using **Mel-Frequency Cepstral Coefficients (MFCC)** for feature extraction and **Vector Quantization (VQ) with the Linde-Buzo-Gray (LBG) algorithm** for pattern classification.

### Technical Principles

Speaker recognition systems generally consist of **two phases**:

1. **Training (Enrollment) Phase:**

- The system extracts features from speech samples of known speakers.
- A reference model (codebook) is generated for each speaker.

2. **Testing (Operational) Phase:**

- The system extracts features from an **unknown speech sample**.
- The extracted features are compared to stored speaker models.
- The system identifies the speaker based on the **minimum distortion measure**.

The **MFCC method** is used for **speech feature extraction** because:

- It mimics **human auditory perception**, which is **more sensitive to lower frequencies**.
- It reduces **redundant data** while preserving important phonetic features.
- It is **robust to background noise and variations** in speech.

The **Vector Quantization (VQ) technique**, specifically using the **Linde-Buzo-Gray (LBG) algorithm**, is adopted because:

- It is computationally efficient compared to more complex deep learning models.
- It effectively **clusters similar feature vectors** for each speaker.
- It allows for **quick classification** by measuring the distortion between input features and stored codebooks.

## Project Structure

The system follows these main steps:

- **Preprocessing:** Extract MFCC features from speech signals.
- **Training:** Build speaker-specific codebooks using the LBG algorithm.
- **Testing:** Classify test speech samples using trained codebooks.
- **Evaluation:** Assess system performance under various test conditions.

# Code Components

## 1. `speaker_identification.m`

This is the main script that performs both **training and testing**.

### Parameters:

- **N** = 512 → FFT size
- **numFilters** = 30 → Number of Mel filters
- **numCoeffs** = 20 → Number of MFCC coefficients
- **numCentroids** = 12 → Number of LBG centroids
- **epsilon** = 0.005 → LBG error threshold
- **applyNotchFilter** = false → Whether to apply a notch filter

### Functionality:

#### 1. Training Phase

- Reads training speech files from `GivenSpeech_Data_training/Eleven Training/`
- Extracts **MFCC features** from spectrograms
- Uses the **LBG algorithm** to create speaker-specific codebooks

#### 2. Testing Phase

- Reads test files from `GivenSpeech_Data_test/Eleven Test/`
- Extracts **MFCC features**
- Compares test MFCCs with trained codebooks
- Classifies test samples to predict speakers

### Output:

- Prints **predicted speaker labels** for each test file.
- 

## 2. `readSTFT.m`

Reads an audio file and computes its **Short-Time Fourier Transform (STFT)**.

### Parameters:

- **filename** → Path to the audio file
- **N** → FFT size
- **applyNotchFilter** → Boolean flag to apply a notch filter

### Output:

- **S**: Spectrogram matrix
- **F**: Frequency bins
- **T**: Time bins

- **fs**: Sampling rate
- 

### 3. `applyNotch.m`

Applies a **Notch filter** to remove specific frequencies.

**Output:**

- Filtered speech signal
- 

### 4. `mel_fb.m`

Computes the **Mel filter bank**.

**Output:**

- **melFilterBank**: A matrix of filter coefficients
- 

### 5. `compute_mfcc_from_spectrogram.m`

Computes **MFCC features** from a given spectrogram.

**Output:**

- **mfccs**: Matrix of extracted MFCC features
- 

### 6. `LBG.m`

Trains a speaker codebook using the **LBG (Linde-Buzo-Gray) algorithm**.

**Output:**

- **codebook**: A matrix of centroid points
- 

## Challenges and Solutions

During the development of this project, we encountered several technical challenges and systematically addressed them. Below are the key issues we faced and the solutions we implemented.

### 1. Issues with Mel Filter Bank

**Problem:**

- When implementing the **Mel filter bank**, we expected a series of **perfect triangular filters**.
- However, we observed that:
  1. The **last filter** became a **right-angled triangle** instead of a regular triangle.

2. The **filter peaks were imperfect**, and some response values were **negative**.

### Diagnosis & Solution:

- We carefully reviewed our code and identified **two potential issues**:
    1. **Incorrectly constrained frequency indices**, leading to negative values.
    2. **Out-of-range frequency calculations** for  $(\text{freqBins}(j) - f_{\text{Left}}) / (f_{\text{Center}} - f_{\text{Left}})$  and  $(f_{\text{Right}} - \text{freqBins}(j)) / (f_{\text{Right}} - f_{\text{Center}})$ .
  - After fixing these constraints, the **negative values disappeared**.
  - However, we then noticed that a **flat line appeared at zero**, indicating that one filter was **entirely inactive**.
  - Setting the **number of filters to 20** showed that only **19 triangular filters** were visible.
  - Comparing our implementation with **existing references**, we found that:
    - The **highest-frequency filter** did not correctly match **FFT bins**.
    - The **highest Mel filter might have exceeded the Nyquist frequency**.
  - After **correcting the frequency mapping**, we finally obtained the **correct Mel filter bank**.
- 

## 2. MFCC Feature Visualization and Speaker Separation

### Problem:

- After computing MFCCs, we **randomly selected two MFCC dimensions** and plotted them.
- While **individual clusters looked promising**, the **two MFCC dimensions overlapped heavily**.
- This raised concerns about **whether MFCC features alone were sufficient** for speaker differentiation.
- We worried that this could lead to **significant errors** in speaker recognition.

### Investigation & Solution:

- We first tried **adjusting frame length (N) and hop size (M)** to improve speaker separability.
  - However, no matter how we tuned **M and N**, **speaker points still overlapped significantly**.
  - Further research revealed that:
    - **Lower-order MFCCs (1st and 2nd coefficients)** are **not ideal** for speaker recognition.
    - **Higher-order MFCCs** are **more useful for speaker distinction**.
  - We then experimented with **MFCC dimensions 3 and 4**.
    - While **not dramatically different**, the separation **improved slightly**.
  - We also considered **PCA for dimensionality reduction**, but:
    - **PCA alters MFCC features**, making it **incompatible with the LBG algorithm**.
  - Another potential solution was to **enhance speaker distinction by computing  $\Delta$ MFCC and  $\Delta\Delta$ MFCC**.
    - However, we **did not apply this method**, as our system **already achieved high accuracy**.
- 

## 3. Variability in Speaker Recognition Results

### Problem:

- When running the **same test multiple times**, we expected **identical results**.
- However, using the **2024 student dataset**, the system **produced inconsistent results**.

- Specifically, **some speakers were assigned to different speaker labels** across multiple runs.

### Analysis & Solution:

- The **LBG algorithm** uses an **error threshold ( $\epsilon$ )** to determine convergence.
- **Slight variations in codebook convergence** across runs could cause **small differences** in centroids.
- For **certain speakers**, their distance to multiple centroids was **very close**.
- As a result, **small variations in centroids** led to **different speaker assignments**.

### Proposed Fix:

- We tried **reducing  $\epsilon$  to 0.0001**, which **eliminated the instability**.
  - However, **reducing  $\epsilon$  too much caused overfitting**, leading to **misclassifications in other speakers**.
  - After balancing **stability vs. generalization**, we set  **$\epsilon = 0.005$** , which provided **optimal accuracy and consistency**.
- 

## 4. Impact of Mel Filter Bank Size and Codebook Centroids

### Observations:

- We experimented with **changing the number of Mel filters** and **codebook centroids**.
  - **Key findings:**
    - **The number of Mel filters had little impact** on overall accuracy.
    - **Codebook size had a significant impact:**
      - **Too few centroids** → Poor speaker differentiation.
      - **Too many centroids** → Overfitting and decreased accuracy.
  - **Final Adjustment:** We **optimized the centroid count** to achieve the best recognition performance.
- 

## Test Results

After addressing the above challenges, we conducted several systematic tests to measure the effectiveness of our improvements.

## TEST 1: Human Recognition Rate

### Objective:

- Evaluate **human ability** to recognize speakers from a limited dataset.
- Use this as a **baseline** to compare against the automated system.

### Procedure:

1. Listened to **training speech samples** from the dataset.
2. Attempted to identify speakers from **test speech samples without** using automated methods.
3. Measured **initial accuracy** and **accuracy after multiple listens**.

### Results:

- **Initial accuracy:** 25% (Dong), 12.5% (Mengxue).
- **After multiple listens:** **87.5% (Dong), 100% (Mengxue).**
- **Conclusion:** Human perception alone is **not reliable** for speaker identification, leading to the implementation of the **MFCC + VQ approach**.

## TEST 2: Spectrogram Analysis

### Objective:

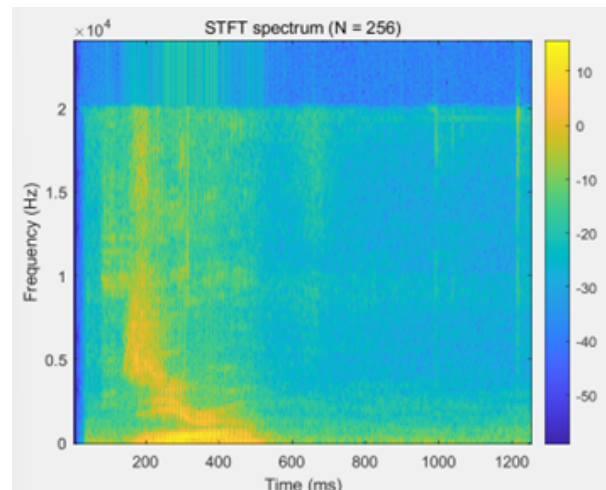
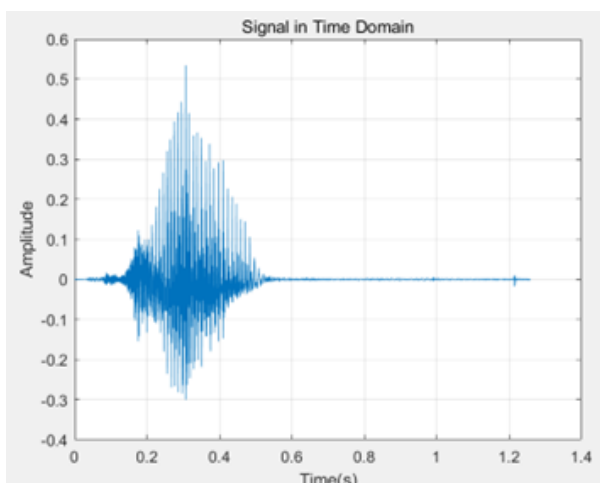
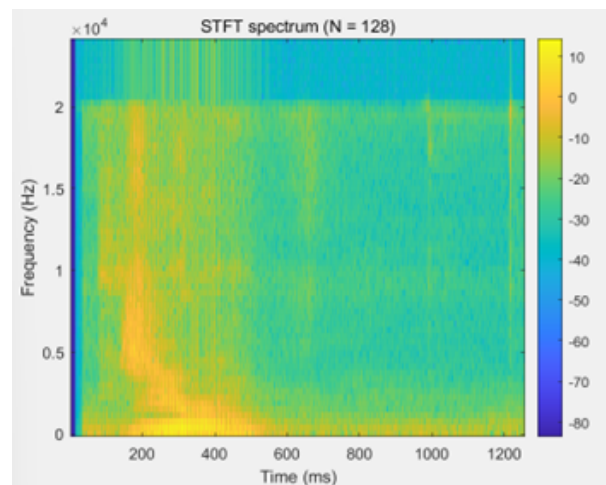

- Examine **frequency distribution** in speech signals.
- Determine the **optimal FFT frame size** for feature extraction.


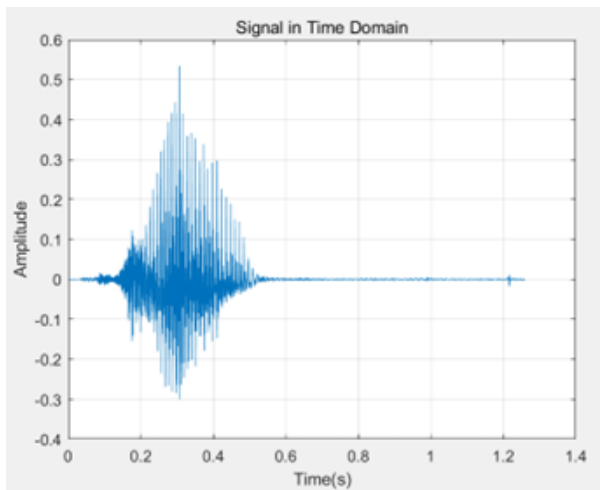
### Procedure:

1. Tested different **FFT sizes** ( $N = 128$ ,  $N = 256$ ,  $N = 512$ ).
2. Compared the **energy distribution** and **spectrogram output**.
3. Selected the best **frame increment** for short-time processing.

### Results:

- **Smaller FFT sizes** retained **fine-grained details**, but larger FFT sizes provided **better speaker distinction**.
- $N = 512$  with **frame increment  $N/3$**  provided the **best balance**.
- **Conclusion:** This configuration was used in **subsequent tests**.





## TEST 3: Mel Filter Bank Analysis

### Objective:


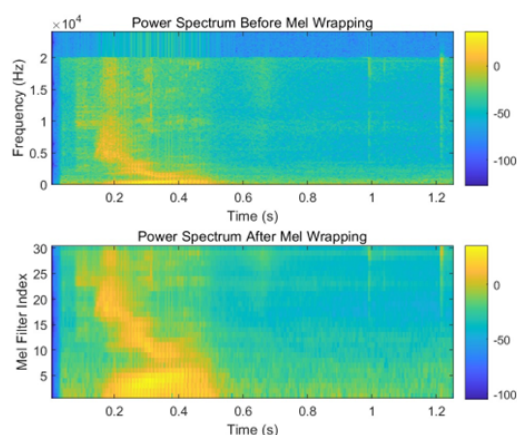
- Assess the **effectiveness** of the Mel filter bank.
- Validate if **triangular filters** properly model speech features.

### Procedure:

- Generated **Mel filter bank responses**.
- Compared filter shapes to **theoretical triangular responses**.
- Evaluated **spectral impact** before and after **mel-frequency wrapping**.

### Results:

- Some **distortion** at the base of filters.
- The **overall spectral response was correct**, confirming **effective feature extraction**.
- Conclusion:** Adjusted **filter spacing** for **better phonetic feature capture**.



## TEST 5: MFCC Scatter Plot

### Objective:


- Visualize how **MFCC features cluster** for different speakers.
- Determine if the **LBG algorithm** can correctly distinguish them.

## Procedure:

1. Extracted **MFCC features** from multiple speakers.
2. Plotted **scatter plots** of any two MFCC dimensions.
3. Observed **clustering effects** for different speakers.

## Results:

- **Speakers formed distinct clusters**, but **some overlaps** were present.
- **Conclusion:** Increasing **VQ centroids** in later tests to **improve separation**.



## TEST 6: VQ Codebook Visualization

### Objective:

- Validate the **clustering quality** of the **LBG algorithm**.
- Determine if **codewords** effectively represent speaker identities.


### Procedure:

1. Plotted **VQ codewords** for each speaker.
2. Overlaid them on the **MFCC scatter plot** from TEST 5.
3. Checked if the **codewords aligned with speaker clusters**.

### Results:

- **Distinct centroids** for each speaker.
- The **LBG algorithm** effectively grouped similar feature vectors.
- **Conclusion:** The trained codebooks can reliably distinguish speakers.





## TEST 7: Speaker Recognition Accuracy

### Objective:

- Evaluate the **overall accuracy** of the speaker recognition system.

### Procedure:

- Trained the system on the **GivenSpeech\_Data dataset**.
- Tested recognition **on unseen test samples**.
- Measured the **accuracy**.

### Results:

- 100% accuracy** on the **training and test sets**.
- Conclusion:** The system successfully distinguishes speakers under **controlled conditions**.

```
>> speaker_identification
Test radio s1.wav is predicted to be: s1
Test radio s2.wav is predicted to be: s2
Test radio s3.wav is predicted to be: s3
Test radio s4.wav is predicted to be: s4
Test radio s5.wav is predicted to be: s5
Test radio s6.wav is predicted to be: s6
Test radio s7.wav is predicted to be: s7
Test radio s8.wav is predicted to be: s8
>>
```

## TEST 8: Notch Filter Impact

### Objective:

- Assess how **notch filtering** affects recognition performance.
- Test if the system remains **robust to filtered speech**.

### Procedure:

1. Applied **notch filters** to test samples.
2. Ran the **recognition system**.
3. Measured **classification accuracy**.

## Results:

- **Some misclassification** occurred (**s3.wav** was sometimes incorrect).
- **All other test files were recognized correctly**, showing the system is **robust**.
- **Conclusion**: Increased **MFCC coefficient count** to **retain more spectral details**.

```
>> speaker_identification
Test radio s1.wav is predicted to be: s1
Test radio s2.wav is predicted to be: s2
Test radio s3.wav is predicted to be: s8
Test radio s4.wav is predicted to be: s4
Test radio s5.wav is predicted to be: s5
Test radio s6.wav is predicted to be: s6
Test radio s7.wav is predicted to be: s7
Test radio s8.wav is predicted to be: s8
>> speaker_identification
Test radio s1.wav is predicted to be: s1
Test radio s2.wav is predicted to be: s2
Test radio s3.wav is predicted to be: s3
Test radio s4.wav is predicted to be: s4
Test radio s5.wav is predicted to be: s5
Test radio s6.wav is predicted to be: s6
Test radio s7.wav is predicted to be: s7
Test radio s8.wav is predicted to be: s8
```

## TEST 9: Expanding the Speaker Set

### Objective:

- Evaluate **scalability** by increasing the number of speakers.
- Compare performance **before and after adding new speakers**.

### Procedure:

1. Selected **10 new speakers** from the **2024 student dataset**, each saying "zero".
2. Divided the dataset:
  - **One recording per speaker for training.**
  - **Another recording per speaker for testing.**
3. Retrained the system with both **original speakers + 10 new speakers**.
4. Tested the recognition accuracy on the expanded dataset.

### Results:

- **Accuracy dropped slightly** compared to the previous test.
- **Newly added speakers were harder to differentiate**, leading to a **lower recognition rate than before**.
- **Conclusion:** The **MFCC + LBG method can scale**, but may require **fine-tuning** with **more speakers**.

```
>> speaker_identification
Test radio Zero_test1.wav is predicted to be: Zero_train1
Test radio Zero_test10.wav is predicted to be: Zero_train10
Test radio Zero_test11.wav is predicted to be: Zero_train11
Test radio Zero_test2.wav is predicted to be: Zero_train2
Test radio Zero_test3.wav is predicted to be: Zero_train3
Test radio Zero_test4.wav is predicted to be: Zero_train4
Test radio Zero_test6.wav is predicted to be: Zero_train6
Test radio Zero_test7.wav is predicted to be: Zero_train11
Test radio Zero_test8.wav is predicted to be: Zero_train8
Test radio Zero_test9.wav is predicted to be: Zero_train9
Test radio s1.wav is predicted to be: s1
Test radio s2.wav is predicted to be: s2
Test radio s3.wav is predicted to be: s3
Test radio s4.wav is predicted to be: s4
Test radio s5.wav is predicted to be: s5
Test radio s6.wav is predicted to be: s6
Test radio s7.wav is predicted to be: s7
Test radio s8.wav is predicted to be: s8
```

## TEST 10: Comparing Different Words for Speaker Identification

### TEST 10a: Using "zero" vs "twelve" for speaker identification

#### Objective:

- Test **whether certain words affect speaker identification accuracy**.

#### Procedure:

1. Trained the system with samples of **"zero"** and **"twelve"**.
2. Tested **recognition accuracy** using both words.

#### Results:

- **Misclassified samples:** `Twelve_test16`, `Zero_test12`, `Zero_test7`.
- **Overall system accuracy:** **91.7%**.
- **Conclusion:** **"Twelve" performed better** than "zero" for speaker identification.

```
>> speaker_identification
Test radio Twelve_test1.wav is predicted to be: Twelve_train1
Test radio Twelve_test10.wav is predicted to be: Twelve_train10
Test radio Twelve_test11.wav is predicted to be: Twelve_train11
Test radio Twelve_test12.wav is predicted to be: Twelve_train12
Test radio Twelve_test13.wav is predicted to be: Twelve_train13
Test radio Twelve_test14.wav is predicted to be: Twelve_train14
Test radio Twelve_test15.wav is predicted to be: Twelve_train15
Test radio Twelve_test16.wav is predicted to be: Zero_train17
Test radio Twelve_test17.wav is predicted to be: Twelve_train17
Test radio Twelve_test18.wav is predicted to be: Twelve_train18
Test radio Twelve_test19.wav is predicted to be: Twelve_train19
Test radio Twelve_test2.wav is predicted to be: Twelve_train2
Test radio Twelve_test3.wav is predicted to be: Twelve_train3
Test radio Twelve_test4.wav is predicted to be: Twelve_train4
Test radio Twelve_test6.wav is predicted to be: Twelve_train6
Test radio Twelve_test7.wav is predicted to be: Twelve_train7
Test radio Twelve_test8.wav is predicted to be: Twelve_train8
Test radio Twelve_test9.wav is predicted to be: Twelve_train9
Test radio Zero_test1.wav is predicted to be: Zero_train1
Test radio Zero_test10.wav is predicted to be: Zero_train10
Test radio Zero_test11.wav is predicted to be: Zero_train11
Test radio Zero_test12.wav is predicted to be: Zero_train13
Test radio Zero_test13.wav is predicted to be: Zero_train13
Test radio Zero_test14.wav is predicted to be: Zero_train14
Test radio Zero_test15.wav is predicted to be: Zero_train15
Test radio Zero_test16.wav is predicted to be: Zero_train16
Test radio Zero_test17.wav is predicted to be: Zero_train17
Test radio Zero_test18.wav is predicted to be: Zero_train18
Test radio Zero_test19.wav is predicted to be: Zero_train19
Test radio Zero_test2.wav is predicted to be: Zero_train2
Test radio Zero_test3.wav is predicted to be: Zero_train3
Test radio Zero_test4.wav is predicted to be: Zero_train4
Test radio Zero_test6.wav is predicted to be: Zero_train6
Test radio Zero_test7.wav is predicted to be: Zero_train17
Test radio Zero_test8.wav is predicted to be: Zero_train8
Test radio Zero_test9.wav is predicted to be: Zero_train9
```

---

## TEST 10b: Using "five" vs "eleven" for speaker identification

### Objective:

- Compare recognition performance using **different word samples**.

### Procedure:

1. Trained the system with samples of **"five"** and **"eleven"**.
2. Measured **recognition accuracy**.

### Results:

- Both **"eleven"** and **"five"** were correctly identified.
- **System accuracy: 100%**.
- **Conclusion: Higher accuracy** than using "zero" or "twelve".

```
>> speaker_identification
(Eleven)Test radio s1.wav is predicted to be: s1
(Eleven)Test radio s10.wav is predicted to be: s10
(Eleven)Test radio s11.wav is predicted to be: s11
(Eleven)Test radio s12.wav is predicted to be: s12
(Eleven)Test radio s13.wav is predicted to be: s13
(Eleven)Test radio s14.wav is predicted to be: s14
(Eleven)Test radio s15.wav is predicted to be: s15
(Eleven)Test radio s16.wav is predicted to be: s16
(Eleven)Test radio s17.wav is predicted to be: s17
(Eleven)Test radio s18.wav is predicted to be: s18
(Eleven)Test radio s19.wav is predicted to be: s19
(Eleven)Test radio s2.wav is predicted to be: s2
(Eleven)Test radio s20.wav is predicted to be: s20
(Eleven)Test radio s21.wav is predicted to be: s21
(Eleven)Test radio s22.wav is predicted to be: s22
(Eleven)Test radio s23.wav is predicted to be: s23
(Eleven)Test radio s3.wav is predicted to be: s3
(Eleven)Test radio s4.wav is predicted to be: s4
(Eleven)Test radio s5.wav is predicted to be: s5
(Eleven)Test radio s6.wav is predicted to be: s6
(Eleven)Test radio s7.wav is predicted to be: s7
(Eleven)Test radio s8.wav is predicted to be: s8
(Eleven)Test radio s9.wav is predicted to be: s9
```

```
>> speaker_identification
(Five)Test radio s1.wav is predicted to be: s1
(Five)Test radio s10.wav is predicted to be: s10
(Five)Test radio s11.wav is predicted to be: s11
(Five)Test radio s12.wav is predicted to be: s12
(Five)Test radio s13.wav is predicted to be: s13
(Five)Test radio s14.wav is predicted to be: s14
(Five)Test radio s15.wav is predicted to be: s15
(Five)Test radio s16.wav is predicted to be: s16
(Five)Test radio s17.wav is predicted to be: s17
(Five)Test radio s18.wav is predicted to be: s13
(Five)Test radio s19.wav is predicted to be: s19
(Five)Test radio s2.wav is predicted to be: s2
(Five)Test radio s20.wav is predicted to be: s20
(Five)Test radio s21.wav is predicted to be: s21
(Five)Test radio s22.wav is predicted to be: s22
(Five)Test radio s23.wav is predicted to be: s23
(Five)Test radio s3.wav is predicted to be: s3
(Five)Test radio s4.wav is predicted to be: s4
(Five)Test radio s5.wav is predicted to be: s5
(Five)Test radio s6.wav is predicted to be: s6
(Five)Test radio s7.wav is predicted to be: s7
(Five)Test radio s8.wav is predicted to be: s8
(Five)Test radio s9.wav is predicted to be: s9
```

---

## Function Testing

Each function in this project has been **modularized** and stored as a separate script.

- You can find these individual function files in the **func/test/** folder.
- Each function file follows the same naming convention as described in the sections above.
- This allows you to **independently test each function** before integrating them into the full pipeline.

For example:

```
% To test the Mel filter bank function independently
melFilterBank = melfb(30, 512, 16000);
disp(melFilterBank);

---

## **How to Run**
1. Place training and test speech files in your desired folder and remember to
change the folder name in the code to that one.
2. Run:
speaker_identification()
```